

Q: Consider a system that applies machine learning techniques to predict fraud transactions:

A: To design a machine learning system for predicting fraud transactions using the provided dataset, the following steps can be taken:

Data Cleaning: Handle missing entries marked by "-". Impute or remove them based on the context or use algorithms that can handle missing data.

Exploratory Data Analysis (EDA): Analyze the dataset to understand patterns and anomalies. Use visualizations to identify trends and outliers.

Feature Selection: Choose relevant features that might indicate fraud. For instance, large transactions or unusual patterns in frequency may be significant.

Model Selection: Choose a suitable model like Random Forest or Gradient Boosting which are effective for classification problems.

Training and Testing: Split the data into a training set and a testing set using stratified sampling to maintain the proportion of fraud cases.

Model Evaluation: Use metrics like precision, recall, and F1-score to evaluate model performance, since the dataset is likely imbalanced.

System Performance Evaluation: Implement cross-validation and adjust model hyperparameters to improve performance.

In Python, pandas can be used for data manipulation, matplotlib for EDA, and scikit-learn for model training and evaluation. Model serialization (like with pickle) is essential for deploying the system.

Q: Consider a list of positive numbers $\{a_1, a_2, \dots, a_n\}$. For any $1 \leq i < j \leq n$, define $V(\text{text}(\text{prod}(a_i, a_{i+1}, \dots, a_j)))$ as the product of numbers in the contiguous subsequence from a_i to a_j . Let $V(f(i))$ be the maximum product of $V(\text{text}(\text{prod}(a_i, a_{i+1}, \dots, a_j)))$ for all contiguous subsequences ending with a_j . Your task is to efficiently find the values of $V(f(1)), V(f(2)), \dots, V(f(n))$ using Bellman's principle of optimality. Provide a Python program that computes these values.

A: To find the values of $f(1), f(2), \dots, f(n)$ efficiently, we can use dynamic programming by applying Bellman's principle of optimality:

Initialize an array F where $F[i]$ represents the maximum product subsequence ending at position i .

Recognize that $F[i]$ can be computed from $F[j-1]$. Specifically, $F[i] = \max(F[j-1] * a[i], a[i])$. This is because the optimal subsequence ending at i is either the optimal subsequence ending at $j-1$ extended by the element $a[i]$, or just the element $a[i]$ itself if the product of the previous subsequence and $a[i]$ is less than $a[i]$.

The base case is $F[1] = a[1]$.

Iterate through the array to fill in all $F[i]$.

```
def max_product_subsequence(arr):
    if not arr: return []
    F = [0] * len(arr)
    F[0] = arr[0]
    for j in range(1, len(arr)):
        F[j] = max(F[j-1] * arr[j], arr[j])
    return F

# Example usage:
arr = [0.5, 1.5, 30, 10, 5, 0.4, 10]
print(max_product_subsequence(arr))
```

Q: Could you please write a Python program that computes the average sentiment scores of tweets related to "Lenovo" and "Dell" for each day from May 7, 2023, to May 16, 2023? These tweets should be obtained from Twitter, and their sentiment analysis should be done using the "flair" tool. Assume that the Bearer Token needed for accessing Twitter is stored in the first line of a file named "BearerToken.txt". Ensure that your program includes all the necessary libraries.

A:

```
import os
import tweepy
import datetime
from flair.models import TextClassifier
from flair.data import Sentence
# Initialize the sentiment analysis model
classifier = TextClassifier.load('en-sentiment')
```

```
def read_bearer_token(file_path):
    with open(file_path, 'r') as file:
        return file.readline().strip()
```

```
def get_tweets(api, query, start_date, end_date):
    tweets = []
    for tweet in tweepy.Cursor(api.search_tweets,
                               q=query,
                               lang='en',
                               since=start_date,
                               until=end_date,
                               tweet_mode='extended').items():
        tweets.append(tweet.full_text)
    return tweets
```

```
def analyze_sentiment(tweets):
    scores = []
    for tweet in tweets:
        sentence = Sentence(tweet)
        classifier.predict(sentence)
        sentiment = sentence.labels[0]
        if sentiment.value == 'POSITIVE':
            scores.append(sentiment.score)
        else:
            scores.append(-sentiment.score)
    return scores
```

```
def main():
    # Load Bearer Token
    bearer_token = read_bearer_token("BearerToken.txt")

    # Setup Twitter API client
    client = tweepy.Client(bearer_token=bearer_token)
```

```
# Define the date range
start_date = datetime.date(2023, 5, 7)
end_date = datetime.date(2023, 5, 16)
```

```
# Iterate over the date range
current_date = start_date
while current_date <= end_date:
    next_date = current_date + datetime.timedelta(days=1)
```

```
# Get tweets for Lenovo
lenovo_tweets = get_tweets(client, "Lenovo", current_date, next_date)
lenovo_scores = analyze_sentiment(lenovo_tweets)
```

```
# Get tweets for Dell
dell_tweets = get_tweets(client, "Dell", current_date, next_date)
dell_scores = analyze_sentiment(dell_tweets)
```

```
# Calculate average sentiment scores
lenovo_avg = sum(lenovo_scores) / len(lenovo_scores) if lenovo_scores
else 0
dell_avg = sum(dell_scores) / len(dell_scores) if dell_scores else 0
```

```
print(f"Date: {current_date} - Lenovo Avg Sentiment: {lenovo_avg}, Dell Avg Sentiment: {dell_avg}")
```

```
# Move to the next date
current_date = next_date
```

```
if __name__ == "__main__":
    main()
```

Q: Write a Python program that takes in three integers 'r', 'c', and 's' (entered one per line), and then constructs two numpy 2D arrays, 'A' and 'B'. Both arrays should have 'r' rows and 'c' columns. Array 'A' should contain a sequence of integers

starting from 's' and incrementing by 1, laid out from left to right and top to bottom. Array 'B' should contain the square roots of the elements in array 'A'. For instance, given the input 3, 4, and 64, array 'A' should be:

```
[[64, 65, 66, 67],
 [68, 69, 70, 71],
 [72, 73, 74, 75]]
```

And array 'B' should be:

```
[[8.0, 8.06, 8.12, 8.19],
 [8.25, 8.31, 8.37, 8.43],
 [8.49, 8.54, 8.60, 8.66]]
```

The program should then print the sum of these two arrays.

A: import numpy as np

```
def create_arrays(r, c, s):
    # Create array A with sequential integers
    A = np.array([s + i for i in range(r * c)]).reshape(r, c)
```

```
# Create array B with square roots of elements in A
B = np.sqrt(A)
```

```
return A, B
```

```
def main():
    # Input handling
    r = int(input("Enter the number of rows (r): "))
    c = int(input("Enter the number of columns (c): "))
    s = int(input("Enter the starting integer (s): "))
```

```
# Create arrays
A, B = create_arrays(r, c, s)
```

```
# Calculate and print the sum of A and B
sum_array = A + B
print("Sum of arrays A and B:\n", sum_array)
```

```
if __name__ == "__main__":
    main()
```

Q: Write a single Python statement that creates and assigns the following DataFrame to the variable 'table':

year	team	wins	draws	losses
2018	HKU	30	6	
2019	HKU	28		
2020	HKU	32		
2018	CU	29		
2019	CU	26		
2020	CU	26		
2018	UST	21		
2019	UST	17	10	11
2020	UST	19	8	11

(c) Use the 'groupby' method in pandas to create a data structure 'teams_wins' that groups the wins by teams. This should result in 'teams_wins.groups' displaying:

```
{'CU': [3, 4, 5], 'HKU': [0, 1, 2], 'UST': [6, 7, 8]}
```

and 'teams_wins.sum()' showing:

team	wins
CU	87
HKU	98
UST	57

(d) Using 'teams_wins', write a single Python statement that displays the team with the largest number of wins, which should appear as:

team	wins
HKU	98

A: Let's break down each part of your question:

(b) Creating the DataFrame

A single Python statement using pandas to create and assign the specified DataFrame to the variable 'table' would look like this:

```
""python
import pandas as pd

table = pd.DataFrame({
    'year': [2018, 2019, 2020, 2018, 2019, 2020, 2018, 2019, 2020],
    'team': ['HKU', 'HKU', 'HKU', 'CU', 'CU', 'CU', 'UST', 'UST', 'UST'],
    'wins': [30, 28, 32, 29, 26, 26, 21, 17, 19],
    'draws': [6, None, None, None, None, None, None, 10, 8],
    'losses': [None, None, None, None, None, None, None, 11, 11]
})
```

(c) Applying the 'groupby' Method

To get the data structure 'teams_wins' which groups the wins in terms of teams, use:

```
""python
teams_wins = table.groupby("team")["wins"]
```

- 'teams_wins.groups' will give {'CU': [3, 4, 5], 'HKU': [0, 1, 2], 'UST': [6, 7, 8]}.

- 'teams_wins.sum()' will give:

```
team
CU      87
HKU     98
UST     57
Name: wins, dtype: int64
```

(d) Team with the Largest Number of Wins

To find the team with the largest number of wins, use:

```
""python
teams_wins.sum().nlargest(1)
```

This will return the team with the highest total wins. Note that the wins for 'HKU' in your example should sum up to 90 (not 98), based on the DataFrame provided.

Please ensure you have pandas installed in your Python environment to execute these statements. You can install it via pip if it's not already installed:

```
""bash
pip install pandas
```

Q: In a lecture, we discussed a program that trains an LSTM model to predict the Hang Seng Index for the next day. Modify this program to train a new LSTM model that predicts the average Hang Seng Index for the following three days. For example, if today is 01/06/2010, the model should predict the average index for the dates 02/06/2010, 03/06/2010, and 04/06/2010.

The initial part of the program is as follows:

```
""python
from datetime import datetime
import pandas_datareader as data

start = datetime.strptime("2004-01-01", "%Y-%m-%d")
end = datetime.strptime("2009-06-30", "%Y-%m-%d")
df = data.DataReader('^HSI', start=start, end=end, data_source='yahoo')
```

```
df = df[['Close']]
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler(feature_range=(0,1))
scaled_data = scaler.fit_transform(np.array(df).reshape(-1,1))

training_size = int(len(scaled_data) * 0.65)
train_data, test_data = scaled_data[:training_size], scaled_data[training_size:]
```

```
from keras.preprocessing.sequence import TimeseriesGenerator
train_generator = TimeseriesGenerator(train_data, train_data, length=100, batch_size=1000)
test_generator = TimeseriesGenerator(test_data, test_data, length=100, batch_size=1000)
```

```
X_train, y_train = list(train_generator)[0][0], list(train_generator)[0][1]
X_test, y_test = list(test_generator)[0][0], list(test_generator)[0][1]
```

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, LSTM
```

```
model = Sequential()
model.add(LSTM(50, return_sequences=True, input_shape=(100,1)))
model.add(LSTM(50, return_sequences=True))
model.add(LSTM(50))
model.add(Dense(1))
model.compile(loss='mean_squared_error', optimizer='adam')
model.fit(X_train, y_train, epochs=5, batch_size=1000, validation_data=(X_test, y_test))
```

```
train_predictions = scaler.inverse_transform(model.predict(X_train))
test_predictions = scaler.inverse_transform(model.predict(X_test))
```

Note: You don't need to import the necessary libraries in your program.

Hint: Here's a snapshot of 'X_train' and 'y_train' constructed by the above program.

A: To modify the LSTM model for predicting the average Hang Seng Index of the following three days, you'll need to adjust the way the target variable 'y_train' and 'y_test' are constructed. The original program predicts the next day's index, but your goal is to predict the average index of the next three days.

Here's a step-by-step guide on how to modify the program:

- "Data Preparation":
 - You need to create a new target variable that contains the average of the next three days' indices for each day in your dataset.
- "TimeseriesGenerator Modification":
 - Modify the 'TimeseriesGenerator' so that 'y_train' and 'y_test' represent the average of the next three days instead of the next day's value.
- "Model Training":
 - The LSTM model architecture and training process remain the same.

Here is the modified section of the code:

```
""python
# ... [previous code for data import and preprocessing]

# Function to calculate the average of the next three days
def create_y(data, n_days=3):
    y = []
    for i in range(len(data) - n_days):
        avg_next_days = np.mean(data[i+1:i+n_days+1])
        y.append(avg_next_days)
    return np.array(y)
```

```
# Create new y for training and testing
y_train_new = create_y(train_period)
y_test_new = create_y(test_period)
```

```
# Update TimeseriesGenerator for new y
train = TimeseriesGenerator(train_period, y_train_new, length=100, batch_size=1000)
test = TimeseriesGenerator(test_period, y_test_new, length=100, batch_size=1000)
```

```
# Get the modified X and y values
X_train, y_train = zip(*[data for data in train])
X_test, y_test = zip(*[data for data in test])
```

```
# X_train and y_train are now tuples of arrays, convert them to arrays
X_train, y_train = np.array(X_train), np.array(y_train)
X_test, y_test = np.array(X_test), np.array(y_test)
```

```
# ... [LSTM model architecture and training]
```

In this modification:

- The 'create_y' function calculates the average of the next three days for each point in the dataset.
- The 'TimeseriesGenerator' objects are updated to use these new 'y_train_new' and 'y_test_new' arrays.
- We then extract the 'X_train', 'y_train', 'X_test', and 'y_test' from these generators and convert them into numpy arrays for training and testing the LSTM model.

Ensure you have the necessary libraries installed and available in your Python environment to run this modified code.

Q: standard option pricing process and codes

A: import numpy as np

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
```

```
# Example data (replace with real data)
data = pd.DataFrame({
    'stock_price': np.random.rand(1000),
    'strike_price': np.random.rand(1000),
    'time_to_expiration': np.random.rand(1000),
    'volatility': np.random.rand(1000),
    'risk_free_rate': np.random.rand(1000),
    'option_price': np.random.rand(1000) # Target variable
})
```

```
# Feature engineering and normalization
features = data.drop('option_price', axis=1)
```

```

target = data['option_price']
scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)

# Splitting the dataset
X_train, X_test, y_train, y_test = train_test_split(features_scaled, target,
test_size=0.2, random_state=42)

# Neural network architecture
model = Sequential([
    Dense(64, input_dim=features.shape[1], activation='relu'),
    Dense(32, activation='relu'),
    Dense(1, activation='linear')
])

# Compile the model
model.compile(optimizer=Adam(lr=0.001), loss='mean_squared_error')

# Train the model
model.fit(X_train, y_train, epochs=100, batch_size=32, validation_split=0.2)

# Evaluate the model
loss = model.evaluate(X_test, y_test)
print('Test Loss:', loss)

```

Q: Researchers are increasingly applying machine learning methods to determine option prices, moving away from traditional models like the Black-Scholes-Merton (BSM) equation. Discuss the limitations of using the BSM equation for option pricing and explain why machine learning might be a more effective approach.

A: The Black-Scholes-Merton (BSM) equation is a foundational model in finance for option pricing, but it has several limitations that have led researchers to explore machine learning methods as alternatives. Here are some key deficiencies of the BSM model and the potential advantages of using machine learning:

Deficiencies of the Black-Scholes-Merton Model

1. **"Assumption of Log-Normal Distribution"**: BSM assumes that the prices of the underlying asset follow a log-normal distribution and exhibit constant volatility. In reality, market volatility is dynamic and can change rapidly due to various factors.
2. **"Constant Interest Rates and Dividends"**: The model assumes constant risk-free interest rates and dividend yields, which is not always the case in real-world markets.
3. **"European Options Only"**: BSM is primarily designed for pricing European options, which can only be exercised at expiration. It doesn't cater well to American options, which can be exercised at any time before expiration.
4. **"No Consideration of Market Sentiment"**: BSM does not account for market sentiment, macroeconomic factors, or other qualitative aspects that can significantly influence option prices.
5. **"Lack of Flexibility"**: The model's rigid assumptions do not adapt well to complex market conditions, exotic options, or instruments with unique features.

Advantages of Machine Learning for Option Pricing

1. **"Handling Non-Linear Relationships"**: Machine learning algorithms, especially deep learning models, are adept at capturing non-linear relationships and complex patterns in data, which are common in financial markets.
2. **"Adapting to Market Conditions"**: Machine learning models can be trained on current market data, allowing them to adapt to changing market conditions and volatility patterns.
3. **"Incorporating Various Data Sources"**: Machine learning can integrate diverse data sources, including market sentiment from news articles and social media, economic indicators, and more, providing a more comprehensive view for pricing.
4. **"Flexibility"**: Machine learning models can be designed to price a wide variety of option types and can be retrained as new data becomes available or market conditions change.
5. **"Potential for Improved Accuracy"**: With sufficient data, machine learning models have the potential to outperform traditional models in terms of prediction accuracy.
6. **"Reducing Assumptions"**: Unlike BSM, machine learning models do not require strict assumptions about market behavior or the distribution of asset prices.

Conclusion

While machine learning offers promising solutions for overcoming the limitations of traditional option pricing models like BSM, it's important to note that these methods come with their own challenges. Machine learning models require large datasets, are computationally intensive, and their "black box" nature can make them difficult to interpret. Additionally, they can be susceptible to overfitting and may require careful tuning and validation to ensure robustness and reliability. Nonetheless, the adaptability and sophistication of machine learning methods make them a compelling area of research in the field of financial modeling and option pricing.